

Rakshith Srinath

rakshithsrinath17@gmail.com | +1-315-621-0543 | LinkedIn | GitHub | Portfolio | Potsdam, New York, USA

Summary

Data Engineer with 3+ years of experience building scalable data pipelines processing 2–3 TB daily, enterprise data warehouses, and production-grade web applications. Master's in Applied Data Science from Clarkson University. Currently serving as Solutions Engineer owning end-to-end product delivery — from data architecture and ETL pipelines to full-stack development and production deployment. Core expertise in Python, SQL, Apache Spark, AWS/Azure, Flask, React, and PostgreSQL, with an active pulse on AI/ML advancements driving the next wave of data engineering.

Technical Skills

Data Engineering:	Apache Spark, Databricks, Delta Lake, Airflow, Kafka, Snowflake, dbt, Hadoop, ETL/ELT, Data Warehousing, Data Modeling, Stream Processing
Languages:	Python, SQL, Java, JavaScript, Bash
Cloud & DevOps:	AWS (S3, EC2, Glue, Redshift, SageMaker), Azure (Databricks, ADF), Docker, Kubernetes, CI/CD, Linux
ML/AI:	Scikit-learn, PyTorch, TensorFlow, LangChain, RAG, MLflow, Feature Engineering, Model Deployment
Web Development:	Flask, FastAPI, React, Bootstrap, HTML/CSS, REST APIs
Databases & Tools:	PostgreSQL, MySQL, SQL Server, MongoDB, Tableau, Grafana, Power BI, Git, Jira

Professional Experience

Solutions Engineer | Life Line Billing Solutions Inc, Texas, USA — Remote Dec 2025 – Present

- Architected end-to-end claims management system replacing manual Excel workflows, designing PostgreSQL schema with dimensional modeling to centralize claims data from third-party integrations
- Built batch ETL pipeline to ingest, validate, and transform claims data from external systems into normalized relational database with audit logging and error handling
- Developed role-based workflow platform with Flask REST APIs and React frontend, enabling 4 user types to track claims lifecycle across configurable stages
- Scaled system beyond pilot to production deployment serving active users, with roadmap to integrate AI-powered automation features

Data Engineer I | LatentView Analytics (Adobe B2B Marketing) — Chennai, India Aug 2021 – Apr 2023

- Built clickstream analytics pipelines processing 2–3 TB daily using PySpark, HiveQL, and SQL Server for Adobe's B2B platform serving 10K+ users
- Developed automated data reconciliation framework using PySpark and Airflow to detect and resolve discrepancies across upstream data sources, reducing manual QA effort by 60%
- Optimized data warehouse with partitioning, bucketing, and columnar formats (Parquet/ORC), achieving 2x query speedup and 50% storage reduction
- Led migration from on-premise Hadoop to Azure Databricks with Delta Lake, reducing ETL runtimes by 40% with 99.9% pipeline uptime

Software Engineer | Skyward Publication — Bengaluru, India May 2020 – Jul 2021

- Engineered ETL pipelines (10K+ records/day) and full-stack applications for small-scale publisher using Python, JavaScript, MySQL, SQL Server, and REST APIs
- Built internal content management dashboard with JavaScript frontend and Python backend, streamlining editorial workflow and reducing publication turnaround time by 25%

Projects

Real-time Streaming Pipeline | Kafka, Spark Streaming, Delta Lake, AWS Oct – Dec 2025

- Built real-time data pipeline processing 100K events/second with exactly-once semantics and sub-3 second latency using Kafka, Spark Streaming, and Delta Lake
- Deployed on AWS with Docker and Kubernetes, implementing auto-scaling based on consumer lag metrics

Retail Analytics Data Platform | Spark, Airflow, Kafka, Snowflake, dbt Jan – Mar 2025

- Built hybrid real-time and batch data platform for retail analytics, ingesting POS transactions via Kafka and processing daily inventory snapshots with Spark and Airflow
- Modeled star schema in Snowflake with dbt transformations, enabling near real-time sales dashboards and inventory forecasting across 500+ SKUs

Enterprise Database Design — Behavioral Data Warehouse | MS Access, VBA, SQL Mar – Dec 2025

- Lead a team of 5 people to architect enterprise data warehouse with dimensional modeling to centralize behavioral data from 50+ distributed entry points across residential facilities
- Built automated ETL pipeline with validation rules, referential integrity constraints, and audit logging, processing 1000+ monthly transactions at 99.9% accuracy

Universal Ergonomics | Python, ChromaDB, FastAPI, React Jan – Present

- Built full-stack e-commerce application with FastAPI backend, React frontend, and user authentication
- As sole founder, managing end-to-end product development including AI-powered inventory management, price prediction features, and continuous feature releases

LLM-Powered Document Analysis System | LangChain, ChromaDB, FastAPI, React May – Jul 2025

- Developed RAG-based document analysis system using LangChain and ChromaDB, achieving 92% retrieval accuracy across 10K+ documents
- Built full-stack application with FastAPI backend, React frontend, user authentication, and real-time document processing

Education

Clarkson University — Potsdam, New York | Master's in Applied Data Science Jan 2024 – Dec 2025

Christ University — Bengaluru, India | Bachelor's of Technology in Computer Science & Engineering Aug 2016 – Jun 2020