

Rakshith Srinath

rakshithsrinath17@gmail.com | +1-315-621-0543 | LinkedIn | GitHub | Potsdam, NY

Summary

Versatile technologist with 3+ years of experience in **data engineering**, **full-stack development**, and **machine learning**. Proven track record building scalable data pipelines processing **2-3 TB daily**, developing web applications, and deploying ML models. Core expertise in **Python**, **SQL**, **Apache Spark**, **cloud platforms (AWS/Azure)**, and modern ML/web frameworks. MS in Applied Data Science candidate at Clarkson University combining engineering depth with analytical rigor.

Technical Skills

Languages & Core: Python, SQL, Apache Spark, Java, JavaScript, Bash, ETL/ELT, Data Warehousing, Data Modeling
Data Engineering: Databricks, Delta Lake, Airflow, Kafka, Snowflake, dbt, Hadoop, Hive, Oozie, Stream Processing
ML/AI: PyTorch, TensorFlow, Scikit-learn, LangChain, RAG, MLflow, Feature Engineering, Model Deployment
Cloud & DevOps: AWS (S3, EC2, SageMaker), Azure (Databricks, ADF), Docker, Kubernetes, CI/CD, Linux
Databases & Tools: PostgreSQL, MySQL, SQL Server, MongoDB, Alteryx, Tableau, Grafana, Power BI, Git, Jira

Professional Experience

Data Engineer I | LatentView Analytics (Adobe B2B Marketing) — Chennai, India Aug 2021 – Apr 2023

- Built and maintained **enterprise-scale clickstream analytics pipelines** processing **2-3 TB daily** using PySpark, HiveQL, and SQL Server for Adobe's B2B marketing platform serving **10K+ users**
- Optimized data warehouse performance through **partitioning, bucketing, and columnar formats** (Parquet/ORC), achieving **2x faster query execution** and **50% storage reduction**
- Spearheaded migration from on-premise Hadoop to **Azure Databricks with Delta Lake**, reducing ETL runtimes by **40%** and improving data freshness from hourly to near real-time
- Implemented **data quality frameworks** and monitoring systems ensuring **99.9% pipeline uptime** and data accuracy

Database Developer | Elmcrest Children's Center — Syracuse, NY (Part-Time) Mar 2025 – Present

- Architected **enterprise data warehouse** with dimensional modeling (fact/dimension tables) to centralize student behavioral data from **50+ distributed entry points** across residential facilities
- Implemented **three-tier architecture**: presentation layer (Access forms), business logic layer (VBA modules), and data layer (normalized relational schema) ensuring scalability and maintainability
- Developed **real-time data pipeline** with automated validation rules, referential integrity constraints, and audit logging, processing **1000+ monthly transactions** with **99.9% accuracy**

Database Developer | Skyward Publication — Bengaluru, India May 2020 – Jul 2021

- Architected **optimized database schemas** in MySQL/SQL Server improving query performance by **30%**, Developed **Python ETL pipelines** processing **100K+ records daily**

Projects

Real-time Streaming Pipeline | Kafka, Spark Streaming, Delta Lake, AWS Oct – Dec 2024

- Built **real-time pipeline** processing **100K events/second** with exactly-once semantics and **<3 second latency**
- Deployed on AWS using Docker/Kubernetes with auto-scaling based on lag metrics

LLM-Powered Document Analysis System | LangChain, ChromaDB, FastAPI, React May – Jul 2024

- Developed **RAG system** achieving **92% accuracy** across **10K+ documents** with web interface
- Built React + FastAPI application with user authentication and real-time document processing

End-to-End ML Pipeline for Churn Prediction | Airflow, MLflow, AWS SageMaker Jan – Mar 2024

- Built production ML pipeline with **feature store** processing **5M customer records daily**
- Deployed XGBoost model with A/B testing, reducing churn by **18%** and saving **\$2M annually**

Apache Spark Benchmarking | Spark, DuckDB, Pandas Feb – Apr 2024

- Benchmarked frameworks on **10M+ flight records** achieving **45% performance improvement**
- Achieved **45% performance improvement** through strategic partitioning, caching, and broadcast joins

System Monitoring Tool | Flask, Docker, Grafana, MySQL Jul 2024 – Present

- Developed containerized **Flask API + Python agent** via Docker Compose to log CPU, memory, and disk metrics in MySQL fact/dim schema

Education

Clarkson University — Potsdam, NY | MS in Applied Data Science Jan 2024 – Dec 2025

Great Lakes Institute of Management — Bengaluru, India | PG Diploma in Data Science Oct 2020 – Jun 2021

Christ University — Bengaluru, India | BS in Computer Science & Engineering Aug 2016 – Jun 2020